



FFRI-NM

FOUNDATION FOR
RELATIONAL INQUIRY

WHEN AI BECOMES EVIDENCE

A Comparative Audit of Machine-Signal Disclosure in University Academic Integrity Procedures

FFRI-NM Project 01 · Research Note 01 · Version 1.0 · 21 September 2026

Research question

When a machine-generated signal enters an academic-misconduct process, how does it acquire institutional force, what part of that signal becomes visible to the affected student, when does disclosure occur, and through what procedural route can the signal itself be contested?

Institutional author: Foundation for Relational Inquiry (FFRI-NM)

Abstract

This research note presents a descriptive comparative audit of public academic-integrity procedures at 24 purposively selected universities in Australia, Canada, the United States and the United Kingdom. It examines how AI-detector outputs move from technical signals into institutional evidence, what part of the signal becomes visible to the affected student, when disclosure occurs, and how the signal can be challenged. The audit separates tool-layer visibility, formal evidence disclosure, human interpretation, signal-level contestability and decision-level review; it also distinguishes multiple disclosure pathways and six timing stages. The corpus documents materially different institutional architectures, including normative exclusion, operational non-use, pre-formal referral, supporting or conditional evidentiary use, local exceptions and central/local tension. An external counterexample shows that direct pre-decision report disclosure is institutionally feasible. The contribution is descriptive and methodological: a source-linked comparative dataset and reproducible codebook, not a prevalence estimate, legal-compliance audit or test of detector accuracy.

Keywords: AI governance; academic integrity; AI detection; evidentiary access; contestability; procedural transparency; institutional policy.

Suggested citation

Foundation for Relational Inquiry (FFRI-NM). (2026). When AI Becomes Evidence: A Comparative Audit of Machine-Signal Disclosure in University Academic Integrity Procedures (Research Note 01, Project 01, version 1.0).

Companion dataset: Foundation for Relational Inquiry (FFRI-NM). (2026). When AI Becomes Evidence: Companion Dataset (Corpus v2.4 / Codebook v0.4).

Publication note: This note reports a descriptive audit of public institutional procedures. It does not rank universities, assess legal compliance, or determine the technical accuracy of AI-detection tools.

© 2026 Foundation for Relational Inquiry (FFRI-NM). Original FFRI-NM text, dataset and codebook: CC BY 4.0. Third-party source materials and the FFRI-NM name/logo are excluded from this license.

24	4	24/24	6
universities	countries	second-pass source re-check	central supporting-signal regimes

Executive summary

This research note examines how university academic-integrity procedures handle machine-generated signals about suspected AI use. The core dataset contains 24 purposively selected universities: 6 in Australia, 6 in Canada, 8 in the United States and 4 in the United Kingdom. Each row was re-checked in a separate second-pass source audit against current official public sources as of 21 September 2026. The study does not test detector accuracy and does not score institutional fairness. Instead, it traces the institutional transition from machine output to formal evidence.

The source-reverified corpus identifies seven distinct architectures. Eight institutions normatively exclude detector output from an evidentiary role; six have disabled or do not operate an institutional detector in the coded scope; six allow a central supporting or conditional machine signal; and four fall into smaller categories involving central/local tension, pre-formal referral, local exception, or no sufficiently explicit institution-wide rule. The only material architecture correction during the second-pass audit was the University of Edinburgh, whose current staff guidance expressly tells staff not to use AI-detection tools for academic-integrity decisions.[1]

The main result is that machine-signal visibility is layered, multi-path and time-dependent rather than binary. Tool-layer visibility, formal evidence disclosure, human interpretation, signal-level challenge and decision-level review are distinct. A report can be hidden at the detector interface yet become accessible later because a general evidence rule applies; conversely, prior notice that a detector may be used does not necessarily give the student access to the case-specific output.

The revised codebook therefore separates (i) F1 cases, where a machine signal may shape suspicion or referral but does not enter formal adjudicative evidence, from F2 cases, where the signal may enter the formal record as supporting or conditional evidence; (ii) multiple disclosure pathways that can coexist; (iii) Signal Access from Signal Challenge Channel; and (iv) six timing stages from pre-use notice to appeal/on-request disclosure.

A separate external codebook review applied the definitions to a stratified eight-institution subset and reported descriptive agreement with the published codes across six selected variables. Because that reviewer had access to the original corpus and did not conduct blind fresh primary-source coding, the exercise is treated as a codebook stress-test rather than formal inter-coder reliability; no Cohen's kappa is reported.

An external counterexample also constrains the substantive claim. University College Birmingham explicitly requires AI-detection reports to be shared with the student in relevant cases and included in the meeting invitation.[2] Direct pre-decision disclosure is therefore institutionally feasible. The defensible contribution of Project 01 is narrower: a reproducible, source-linked comparative audit of how published procedures move a machine signal from tool output into a record that can be inspected, interpreted, challenged, or excluded.

1. Question and scope

Universities increasingly face allegations that student work contains unauthorized generative-AI output. A detector score may be one input into that process, but the consequential institutional act is not the score itself. The key transition occurs when a machine-generated signal is taken up by a human process: it may prompt attention, trigger an investigation, enter a formal evidence record, be interpreted by staff, or be excluded from adjudication altogether.

Project 01 asks what happens at that transition. Specifically, it examines whether detector output has institutional force; whether it enters a formal record; what machine-derived object the student can access; when access occurs; whether the student can challenge the signal itself; and whether access arises from an AI-specific rule, a general evidence rule, staff discretion, supporting-information-only, or no publicly specified path.

This framing deliberately avoids two adjacent questions. First, the project does not determine whether a detector is technically accurate. Recent literature has already examined detector reliability and procedural-fairness concerns.[6] Second, it does not propose a new general theory of contestability or evidentiary rights. Prior work already establishes the importance of evidentiary access and the institutional preconditions of challenge.[3][4]

2. Method

2.1 Sampling and source-retrieval protocol

The corpus is purposive rather than representative. It contains 24 universities selected through maximum-variation sampling to expose contrasting institutional approaches across Australia, Canada, the United States and the United

Kingdom. The unit of analysis is the published institutional procedure, not an individual misconduct case. Country allocation was not quota-based; institutions were selected for architectural variation rather than national balance.

An initial scan of official university materials identified contrasting detector-policy states, including normative exclusion, operational non-use, central conditional or supporting use, local variation, and unresolved central policy. Institutions were added to ensure that these architectures were represented. Selection was not based on expected disclosure outcomes or on a judgment that an institution was procedurally "strong" or "weak".

Inclusion required: (1) official public institutional material concerning AI detection, a documented non-use or exclusion position, or an explicitly bounded local rule; (2) a public academic-integrity or misconduct procedure with identifiable scope; and (3) enough source material to code the machine-signal pathway without silently generalizing a local rule to the whole institution. The core corpus was closed at 24 after the separate second-pass source audit stabilized the descriptive architecture. Later searches were used as adversarial stress tests rather than added to prevalence counts.

Sources were checked against current official public pages as of 21 September 2026. The companion workbook records official URLs, institutional scope, source role, audit status and retrieval date. Because public web pages can change, the corpus describes published institutional architecture at the audit date rather than a permanent policy state.

2.2 Layered coding model and Codebook v0.4

The audit separates four institutional stages:

A / Tool-Interface	->	B / Formal Evidence	->	C / Interpretation-Response	->	D / Decision-Review
---------------------------	----	----------------------------	----	------------------------------------	----	----------------------------

The v0.4 codebook makes the F1/F2 boundary explicit. F1 is used when a detector signal may influence suspicion, referral, human attention, investigation, or later human testimony, but the machine output itself does not enter formal adjudicative evidence. F2 is used when the machine output may enter the formal record as supporting or conditional evidence, even if sole reliance is prohibited. Human evidence prompted by a detector does not by itself convert the underlying detector output from F1 to F2.

Disclosure Path is no longer a single-choice field. A case can carry more than one pathway, such as Melbourne, where discretionary early PDF sharing can coexist with later access through a general evidence-review rule. The public dataset therefore records separate flags for direct AI-specific entitlement, generic evidence entitlement, staff discretion, supporting-information-only, no adjudicative force, and other/not found.

Signal contestability is split into two variables. Signal Access asks whether the affected student can obtain the case-specific machine-derived object. Signal Challenge Channel asks whether the student can challenge that machine-derived signal itself. A general appeal against the final institutional decision does not automatically satisfy either variable.

Human interpretation is coded separately through three variables. HI1 (Human Evaluation Required) asks whether a human actor must evaluate, validate, reject, or contextualize the machine signal before consequential action. HI2 (Interpretation Communicated) asks whether the affected student receives an intelligible human explanation of why the signal is institutionally relevant. HI3 (Interpretation Contestable) asks whether the student can answer or challenge that human interpretation. These variables distinguish access to a machine output from access to the institutional meaning attached to it.

2.3 Timing taxonomy

Disclosure timing is coded independently of disclosure path:

Stage	Definition	Interpretive caution
T0	Before machine use	Notice that a tool may be used is not the same as access to the resulting signal.
T1	After signal, before formal allegation	Can enable early response before formal case construction.
T2	With formal allegation / initial investigation	Signal or supporting information enters the student-facing formal process.
T3	Before hearing or final decision	Disclosure occurs later, but still before the consequential decision.
T4	With decision	The relevant machine-related information becomes visible only with the outcome.
T5	Appeal / on request	Access arises only through review, appeal

Stage	Definition	Interpretive caution
		or a specific request.

2.4 Second-pass source audit and external codebook review

The second-pass audit preserved the frozen first-pass corpus and re-checked all 24 rows in a separate source-audit pass. One substantive architecture change resulted: Edinburgh moved from "operational absence / disabled" to "normative exclusion," because current staff guidance expressly says not to use AI-detection tools for academic-integrity decisions.[1] Oregon State remains coded as "institution-wide detector rule not found": strong college-level guidance exists, but the project does not generalize a faculty or college rule to the entire university.

A separate external codebook review then applied the published definitions to a stratified subset of eight institutions (two selected from each country). The reviewer reported descriptive agreement with the original coding across six selected variables in all eight cases. The review also identified the F1/F2 boundary, multi-path disclosure, signal-access versus signal-challenge separation, FOC discipline, timing, human interpretation, and rule discoverability as areas requiring clarification. Those clarifications are incorporated in Codebook v0.4.

This exercise is not formal inter-coder reliability. The reviewer had seen the original corpus, coded from the Research Note and companion dataset rather than conducting blind fresh primary-source coding, and worked with a small n. Accordingly, the project reports descriptive agreement only and does not report Cohen's kappa or Krippendorff's alpha.

3. Results

3.1 Institutional architectures

The source-reverified 24-row corpus produces the following descriptive architecture:

Architecture	Count	Share	Meaning
F0N - Normative exclusion	8	33.3%	Detector output excluded from the formal evidentiary role.
F0O - Operational absence / disabled	6	25.0%	No active institutional detector signal in the coded scope.
F0C - Central exclusion with local/legacy tension	1	4.2%	Central guidance gives detector output no evidentiary force, but materially relevant local/legacy variation remains.
F1 - Pre-formal attention / referral	1	4.2%	A machine score may shape concern or referral but does not enter formal adjudicative evidence.
F2 - Central supporting / conditional signal	6	25.0%	Detector output may enter the central process as supporting or conditional evidence.
F3 - Local-exception architecture	1	4.2%	Documented local use may occur without a central detector-evidence standard.
F4 - Central rule not found	1	4.2%	No sufficiently explicit institution-wide detector rule located in reviewed public material.

Fourteen of the 24 institutions therefore either normatively exclude detector output from the evidentiary role or do not operate an institutional detector in the coded scope. Six centrally permit a supporting or conditional machine signal. The remaining four cases are kept separate because they represent qualitatively different policy structures rather than intermediate points on a single scale. These counts describe the purposive corpus only; they are not prevalence estimates.

3.2 The six central supporting-signal regimes

The six F2 institutions are the most informative for the disclosure question because the machine signal can acquire formal institutional force. Their public procedures do not converge on a single pathway.

Institution	Signal role	Disclosure pathway(s)	Signal access / challenge
University of Missouri	Supporting only; not sole evidence	Generic evidence entitlement if report enters the formal record	Access CONDITIONAL; challenge YES
UNSW Sydney	Investigative flag / trigger	Supporting-information-only; full-	Access PARTIAL; challenge

Institution	Signal role	Disclosure pathway(s)	Signal access / challenge
		report entitlement not found	PARTIAL [7]
University of Chicago	Supporting information / conversation guide	Transparent human-led conversation; full-report entitlement not found	Access PARTIAL; challenge PARTIAL
Michigan State University	Potential indicator / trigger; never sole basis	Prior-use transparency; case-specific report path not found	Access NOT FOUND; signal challenge NOT FOUND
University of Sydney	Supporting evidence; not sole evidence	Generic evidence entitlement when material is used to decide the case	Access CONDITIONAL; challenge YES [8][9]
University of Melbourne	Indicator / supporting evidence; second item required	Dual path: staff discretion + generic evidence review	Access CONDITIONAL; challenge PARTIAL [10][11]

3.3 Timing is an independent procedural dimension

The revised timing codes show that "disclosure" does not identify a single procedural moment. Within the six F2 regimes, Sydney and UNSW place relevant information at T2, with the allegation or at the start of formal investigation. Missouri and the University of Chicago place material access or explanation later at T3, before a hearing or final conclusion. Melbourne can operate through two stages: discretionary early sharing after concern (T1) and formal evidence review before decision (T3). Michigan State provides T0 policy-level notice about detector use when an instructor chooses to use one, but the public materials reviewed do not establish when the student can access the case-specific machine output.

The external UCB counterexample is a T2 architecture: the AI-detection report accompanies the meeting invitation and is therefore available before the student responds.[2] The practical implication is descriptive rather than normative: two institutions can both "disclose" machine-related information while giving the affected person materially different opportunities to inspect and answer it before institutional judgment.

Institution	Status	Earliest relevant stage	Note
Missouri	Core F2	T3	Evidence access during investigation / before hearing.
UNSW Sydney	Core F2	T2	Relevant supporting information with formal allegation / during investigation.
University of Chicago	Core F2	T3	Human-led conversation before conclusion; full-report right not explicit.
Michigan State	Core F2	T0 for notice; case-specific timing NOT FOUND	Prior-use transparency does not establish output access.
University of Sydney	Core F2	T2	Supporting material with written allegation / before response.
University of Melbourne	Core F2	T1 + T3	Discretionary early PDF sharing plus later formal evidence review.
University College Birmingham	External comparator	T2	AI report accompanies meeting invitation / before response.

3.4 Applicable-rule discoverability

A second refinement concerns whether an affected student can determine which rule applies to the relevant academic unit. This is not a fairness score. It records the structure of the published rule environment using explicit coding criteria.

CLEAR CENTRAL is used when a central rule, or directly linked central rule set, identifies the applicable detector position without requiring resolution of a source conflict. MULTI-DOCUMENT RECONSTRUCTABLE is used when two or more compatible official documents are needed but together yield one unambiguous pathway. PARTIAL / LOCAL VARIATION is used when a central position exists but applicability materially depends on faculty, department or programme rules. CENTRAL-LOCAL TENSION requires a clear central position plus documented local or legacy material pointing in a materially different direction. CENTRAL RULE NOT FOUND is reserved for cases in which no sufficiently explicit central detector rule was located after the defined search.

In the 24-institution corpus, 15 cases present a clear central detector position, 6 require reconstruction across multiple official pages, 1 (Alberta) preserves documented local variation, 1 (Cambridge) presents a material central/local tension, and 1 (Oregon State) lacks a sufficiently explicit institution-wide detector rule in the reviewed public material. Cambridge and Alberta therefore remain analytically useful not because the audit resolves their local practice, but because they show that policy scope can itself become part of the information problem faced by an affected person.

4. Main finding: disclosure is layered, multi-path and time-dependent

The first-pass variable "Signal_Disclosure" was too coarse. It treated visibility as if a machine report were simply visible or invisible. The revised corpus separates at least five questions: is the report visible in the technical interface; can it enter the formal evidence record; which procedural pathway makes any machine-derived object accessible; when disclosure occurs; and can the student challenge the machine signal itself?

This matters because a report can be technically invisible to a student but procedurally accessible later. Melbourne is a clear example: Turnitin does not make the AI indicator or report visible to students, staff can share a PDF, and the formal misconduct process separately requires human evaluation and evidence review.[10][11] UNSW likewise states that the report is visible to teachers, not students, while treating the detector result only as a flag for further investigation.[7]

The converse also holds. Prior notice that a detector may be used does not itself create access to the case-specific output. Nor does a strong appeal process prove that the person could inspect or challenge the machine signal at the point it first influenced the institution. Signal access, signal challenge and decision-level review are therefore coded as separate variables.

The main descriptive result is not that universities hide machine signals, nor that they disclose them. It is that institutional visibility is staged and procedurally constructed. The same machine output can have different practical meaning depending on whether it remains an internal alert, enters the formal record, is summarized rather than disclosed, is shared only at staff discretion, or becomes evidence that the affected person can inspect and challenge before a decision.

5. Stress test: a direct-disclosure counterexample

Project 01 deliberately searched for counterexamples to the early hypothesis that machine-signal access begins only after suspicion has already been formed. University College Birmingham falsifies that universal claim. Its current Code of Practice states that high Turnitin Similarity/AI detection scores are not definitive evidence, that AI detection reports should be shared with the student, and that any relevant report should accompany the invitation to the meeting.[2]

This counterexample is methodologically important. It shows that full pre-decision report disclosure is not technically or procedurally impossible. It also prevents the project from converting a pattern in a purposive 24-institution sample into a general claim about universities. UCB is used as a positive architecture comparator, not as evidence about prevalence.

6. Relation to prior work and novelty discipline

The project sits inside an already developed literature on contestability, evidentiary access and academic-integrity evidence. Stewart argues that explanation without access to evidence is insufficient for meaningful contestation and develops a broader evidentiary-rights framework.[3] Kahl develops "capture of answerability" in AI-mediated institutional decision-making and explicitly applies the framework to university AI-assisted academic-integrity processes.[4] These works preclude any credible claim that Project 01 introduces evidentiary access, institutional answerability or contestability as new concepts.

The empirical literature is also advancing rapidly. Muñoz and colleagues analyzed 1,162 GenAI-related academic-misconduct cases and 1,855 evidence items from one Australian university, producing a 15-code evidence taxonomy and finding that detection-tool outputs had weak inferential-force profiles relative to several other evidence types.[5] Bassett and colleagues separately argue against educational reliance on AI detectors on methodological and procedural-fairness grounds.[6] Project 01 therefore does not attempt to re-run the detector-accuracy or detector-fairness debate.

A further nearby resource is the University AI Policy Tracker, which indexes hundreds of university records and source-backed claims concerning AI detectors.[12] The tracker is valuable for discovery and validation, but does not itself reconstruct all of the staged pathways examined here. The narrower contribution pursued by Project 01 is an open, reproducible coding of the transition from machine output to formal evidence, including institutional force, disclosure pathway, timing, signal access, signal challenge and applicable-rule discoverability.

Disclosure should not be read as validation of the underlying detector. A procedurally contestable signal can still be technically weak, biased or inferentially inappropriate. Those properties are outside the design of this audit and remain separate empirical questions.

7. Implications for practice

- Knowing that AI detection may be used is not equivalent to seeing the case-specific machine output.

- Tool-interface visibility should not be used as a proxy for procedural evidence access; a report hidden in the tool may later become disclosable through a general evidence rule.
- Institutions that use detector output may need to distinguish explicitly between AI-specific disclosure, general evidence-disclosure rights, staff discretion, and supporting-information-only pathways.
- Where central and local rules differ, discoverability of the applicable rule becomes a procedural question of its own.
- Disclosure does not validate the detector. A contestable signal may still be technically weak, biased, or inferentially inappropriate; those properties are outside this audit.

8. Limitations

- The sample is purposive and limited to 24 institutions in four countries. Percentages and counts describe this corpus; they are not estimates of prevalence across higher education.
- The study codes published institutional architecture, not observed practice. A rule can exist without being consistently followed, and a practice can occur without being fully documented online.
- NOT FOUND does not mean NO. It means the relevant rule was not located in the reviewed official public materials.
- Local, faculty, school and course rules may differ from central university rules. The audit preserves scope and does not silently generalize local guidance. The new `Applicable_Rule_Discoverability` variable describes, but does not resolve, central/local ambiguity.
- The project does not measure detector accuracy, bias, case outcomes, reversal rates, student trust, workload or causal effects of a disclosure regime. It also does not include interviews with students or staff. Such work would require a separate empirical and ethics design.
- The layered codebook remains a research instrument under development. It is not presented as a universal fairness standard or a legal-compliance test.
- All primary coding and source verification in this release were conducted within one research workflow. The external eight-case review provides a useful descriptive agreement check but is not blind primary-source double-coding. Formal inter-coder reliability is therefore not reported.
- The companion dataset records source URLs and retrieval dates, but this public release does not redistribute archival snapshots of every mutable web page. A fuller snapshot-based provenance package may be added in a future archival release.

9. Conclusion

The first substantive FFRI-NM case study identifies a concrete institutional boundary: the point at which a machine-generated signal moves from an internal technical output into a record that can influence a consequential human process. Public university procedures handle that boundary in materially different ways.

The strongest descriptive result is that visibility is staged. Tool-layer access, formal evidence disclosure, human interpretation, signal-level challenge and decision-level review are distinct. A report can be unavailable at one stage and accessible at another. Prior notice can exist without case-specific access. Formal appeal can exist without proving that the upstream machine output was inspectable. Timing further changes the practical structure of contestation: disclosure before response is not the same architecture as disclosure shortly before a hearing or only through later review.

For FFRI-NM, the practical output is an open comparative dataset and a reproducible layered codebook. The immediate next methodological step is not to enlarge the theory, but to test Codebook v0.4 through genuinely blind primary-source double-coding and to strengthen the provenance package. Separate follow-up work could examine longitudinal policy change, wider national sampling, local/central rule dynamics, and whether the documented pathways correspond to student understanding or observed institutional practice.

Appendix A. Core sample and reverified architecture

All rows below were re-checked in a separate second-pass source audit against current official public sources on 21 September 2026. The complete source register, staged fields, timing codes, human-interpretation variables, discoverability criteria and audit notes are contained in the companion workbook *Corpus v2.4 / Codebook v0.4* (Public Release).

Institution	Country	Architecture
Monash University	Australia	F0N
University of Victoria (UVic)	Canada	F0N
Washington State University	USA	F0O

Institution	Country	Architecture
University of Missouri	USA	F2
UNSW Sydney	Australia	F2
University of Chicago	USA	F2
University of Toronto	Canada	F0N
Vanderbilt University	USA	F1
University of California, Berkeley	USA	F0O
Indiana University Bloomington	USA	F0N
Oregon State University	USA	F4
Michigan State University	USA	F2
University of Sydney	Australia	F2
University of Melbourne	Australia	F2
Australian National University (ANU)	Australia	F0O
University of Queensland (UQ)	Australia	F0O
University of British Columbia (UBC)	Canada	F0O
University of Waterloo	Canada	F0O
McGill University	Canada	F0N
University of Alberta	Canada	F3
University of Oxford	United Kingdom	F0N
University of Cambridge	United Kingdom	F0C
University of Edinburgh	United Kingdom	F0N
University College London (UCL)	United Kingdom	F0N

References and representative sources

The companion dataset contains the complete institution-by-institution source register. The references below support the principal methodological and comparative claims in this note.

- [1] University of Edinburgh. "Using Generative AI in Your Work: Guidance for Staff." Updated 24 March 2026. Current guidance states not to use AI-detection tools for academic-integrity decisions. [Source](#)
- [2] University College Birmingham. Code of Practice on Plagiarism and Academic Misconduct. 2025. Relevant AI-detection reports are to be shared with students and included with the meeting invitation. [Source](#)
- [3] Stewart, M. "Beyond Explanation: Evidentiary Rights for Algorithmic Accountability." FAccT 2026. DOI: 10.1145/3805689.3806532. [Source](#)
- [4] Kahl, P. "The Capture of Answerability: AI Governance and Institutional Accountability." Lex et Ratio Ltd Working Paper, 9 May 2026. DOI: 10.5281/zenodo.20093912. [Source](#)
- [5] Muñoz, A., Hinchcliff, M., Langfield, C., et al. "How strong is the evidence in generative AI-related academic misconduct allegations? A mixed-methods analysis." International Journal for Educational Integrity 22, 26 (2026). DOI: 10.1007/s40979-026-00235-9. [Source](#)
- [6] Bassett, M. A., Bradshaw, W., Bornsztejn, H., et al. "Heads we win, tails you lose: AI detectors in education." Journal of Higher Education Policy and Management (2026). DOI: 10.1080/1360080X.2026.2622146. [Source](#)
- [7] UNSW Sydney. "Access to AI tools." Current staff guidance: the Turnitin AI report is visible to teachers, not students, and functions as a flag for investigation. [Source](#)
- [8] University of Sydney. "Artificial intelligence." Current student guidance on AI use and Turnitin AI detection. [Source](#)
- [9] University of Sydney. "Detection and investigation of academic integrity breaches." Current procedural guidance. [Source](#)
- [10] University of Melbourne. "Turnitin's AI writing detection tool." Staff guidance: AI score is an indicator and a second piece of evidence is required before an allegation can proceed. [Source](#)
- [11] University of Melbourne. "Academic misconduct process." Current procedural guidance on allegation, evidence review and response. [Source](#)
- [12] University AI Policy Tracker. "AI detectors." Source-linked public policy tracker used as a discovery and validation resource. [Source](#)

Companion data artifact: FFRI_NM_Project01_When_AI_Becomes_Evidence_Corpus_v2.4_PUBLIC_RELEASE.xlsx